

# ML Project Code

December 9, 2018

In [65]: *#Import required files*

```
import pandas as pd
from sklearn.model_selection import train_test_split
from scipy.cluster.hierarchy import dendrogram, ward, cut_tree
from scipy.spatial.distance import pdist
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
import scipy.spatial.distance as ssd
from sklearn.preprocessing import OneHotEncoder
from sklearn.linear_model import LinearRegression
from sklearn import linear_model
from sklearn.linear_model import Ridge
from sklearn.linear_model import RidgeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn import metrics
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.preprocessing import PolynomialFeatures
```

In [66]: *#load data from csv*

```
df = pd.read_csv('crimes.csv', sep=',')

columns = ['year', 'month', 'day', 'hour', 'location_type', 'zipcode', 'x_coord', 'y_coord', 'o

#columns = ['year', 'zipcode', 'offence_type', 'x_coord', 'y_coord', 'severity']

# create dataframe with relevant columns

df_2 = df[columns]
```

In [67]: *# Enumerate output column:*

```

crimes = pd.factorize(df_2['offence_type'])
df_2['offence_type'] = crimes[0]
crime_list = crimes[1] #create list of crime references that are coded to specific fact

# Enumerate input variables:

year = pd.factorize(df_2['year'])
df_2['year'] = year[0]
year_list = year[1]

month = pd.factorize(df_2['month'])
df_2['month'] = month[0]
month_list = month[1]

day = pd.factorize(df_2['day'])
df_2['day'] = day[0]
day_list = day[1]

hour = pd.factorize(df_2['hour'])
df_2['hour'] = hour[0]
hour_list = hour[1]

location_type = pd.factorize(df_2['location_type'])
df_2['location_type'] = location_type[0]
location_type_list = location_type[1]

# council_district = pd.factorize(df_2['council_district'])
# df_2['council_district'] = council_district[0]
# council_district_list = council_district[1]

# apd_sector = pd.factorize(df_2['apd_sector'])
# df_2['apd_sector'] = apd_sector[0]
# apd_sector_list = apd_sector[1]

# apd_district = pd.factorize(df_2['apd_district'])
# df_2['apd_district'] = apd_district[0]
# apd_district_list = apd_district[1]

```

/anaconda/envs/py35/lib/python3.5/site-packages/ipykernel\_launcher.py:4: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#>  
after removing the cwd from sys.path.  
/anaconda/envs/py35/lib/python3.5/site-packages/ipykernel\_launcher.py:10: SettingWithCopyWarning  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#>  
# Remove the CWD from sys.path while we load stuff.  
/anaconda/envs/py35/lib/python3.5/site-packages/ipykernel\_launcher.py:14: SettingWithCopyWarning  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#>  
/anaconda/envs/py35/lib/python3.5/site-packages/ipykernel\_launcher.py:18: SettingWithCopyWarning  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#>  
/anaconda/envs/py35/lib/python3.5/site-packages/ipykernel\_launcher.py:22: SettingWithCopyWarning  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#>  
/anaconda/envs/py35/lib/python3.5/site-packages/ipykernel\_launcher.py:26: SettingWithCopyWarning  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#>

```
In [68]: df_2=df_2.dropna(axis=0)
        X = df_2.drop(['offence_type','severity'],axis=1).values #sets x and converts to an array
        print(len(X))
        #print(X)

        y = df_2['severity'].values #sets y and converts to an array

        # Split the data into train and test sets for numeric encoded dataset:

        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.20, random_state=
224499
```

```
In [69]: from sklearn.decomposition import PCA

        pca = PCA(4, svd_solver='randomized').fit(X_train)
        comp = pca.transform(X_train)
        comp_test = pca.transform(X_test)
        # proj = pca.inverse_transform(comp)
```

```
In [70]: #Random forest classifier

        classifier = RandomForestClassifier(n_estimators = 100, criterion = 'entropy', random_state=
```

```

classifier.fit(X_train, y_train)
y_pred = classifier.predict(X_test)

print(accuracy_score(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test,y_pred, target_names=crime_list))

```

0.7418040089086859

[31521	50	18	704	57	18	0]			
[ 2240	32	3	75	13	0	0]			
[ 942	4	10	29	7	5	0]			
[ 4961	3	1	932	22	17	1]			
[ 2218	9	5	126	27	27	0]			
[ 0	0	0	0	0	785	0]			
[ 37	0	0	1	0	0	0]]			
			precision	recall	f1-score	support			
Theft	0.75	0.97	0.85	32368					
Robbery	0.33	0.01	0.03	2363					
Auto Theft	0.27	0.01	0.02	997					
Aggravated Assault	0.50	0.16	0.24	5937					
Burglary	0.21	0.01	0.02	2412					
Rape	0.92	1.00	0.96	785					
Murder	0.00	0.00	0.00	38					
avg / total	0.66	0.74	0.66	44900					

In [7]: *#Random forest classifier PCA*

```

classifier = RandomForestClassifier(n_estimators = 100, criterion = 'entropy', random_st
classifier.fit(comp, y_train)
y_pred = classifier.predict(comp_test)

print(accuracy_score(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test,y_pred, target_names=crime_list))

```

0.7192992182735488

[24207	217	75	1143	244	20	0]			
[ 1821	74	15	97	25	0	1]			
[ 780	21	28	24	16	5	0]			
[ 3621	26	5	916	66	20	0]			
[ 1611	23	26	164	75	27	1]			
[ 0	0	1	3	2	648	0]			
[ 24	1	0	0	1	0	0]]			
			precision	recall	f1-score	support			

Theft	0.75	0.93	0.84	25906
Auto Theft	0.20	0.04	0.06	2033
Aggravated Assault	0.19	0.03	0.05	874
Burglary	0.39	0.20	0.26	4654
Robbery	0.17	0.04	0.06	1927
Rape	0.90	0.99	0.94	654
Murder	0.00	0.00	0.00	26
avg / total	0.63	0.72	0.66	36074

```
In [8]: #decision tree classifier
        kfold = KFold(n_splits=10, shuffle=True, random_state=2)
        for k in range(1,15):
            tree= DecisionTreeClassifier(max_depth=k,random_state=2)
            fold_accuracies = cross_val_score(tree, X_train, y_train, cv=kfold)
            print("k:",k)
            print("Cross-validation score:\n{}".format(fold_accuracies))
            print("average cross_validation score: {:.2f}".format(fold_accuracies.mean()))

k: 1
Cross-validation score:
[0.73742204 0.73284823 0.72619543 0.73534304 0.73562024 0.72959113
 0.73241389 0.73574052 0.73587913 0.73185945]
average cross_validation score: 0.73
k: 2
Cross-validation score:
[0.73742204 0.73284823 0.72619543 0.73527374 0.73562024 0.72959113
 0.73241389 0.73574052 0.73587913 0.73185945]
average cross_validation score: 0.73
k: 3
Cross-validation score:
[0.73749134 0.73284823 0.72619543 0.73527374 0.73562024 0.72966043
 0.73241389 0.73574052 0.73601774 0.73185945]
average cross_validation score: 0.73
k: 4
Cross-validation score:
[0.73749134 0.73284823 0.72619543 0.73527374 0.73562024 0.72959113
 0.73241389 0.73574052 0.73601774 0.73179014]
average cross_validation score: 0.73
k: 5
Cross-validation score:
[0.73832294 0.73284823 0.72619543 0.73527374 0.73562024 0.73014553
 0.73303763 0.73574052 0.73587913 0.7326218 ]
average cross_validation score: 0.73
k: 6
```

Cross-validation score:  
[0.73804574 0.73367983 0.72674983 0.73541234 0.73652114 0.73090783  
0.73276041 0.7368494 0.73643357 0.73276041]  
average cross\_validation score: 0.73  
k: 7  
Cross-validation score:  
[0.73797644 0.73381843 0.72681913 0.73534304 0.73652114 0.73139293  
0.73282972 0.73657218 0.73657218 0.73289902]  
average cross\_validation score: 0.73  
k: 8  
Cross-validation score:  
[0.73887734 0.73451143 0.72875953 0.73707554 0.73742204 0.73173943  
0.73622566 0.73788897 0.73795828 0.73671079]  
average cross\_validation score: 0.74  
k: 9  
Cross-validation score:  
[0.73970894 0.73575884 0.72896743 0.73742204 0.73728344 0.73409563  
0.73574052 0.73816619 0.73892855 0.73740384]  
average cross\_validation score: 0.74  
k: 10  
Cross-validation score:  
[0.73846154 0.73562024 0.72875953 0.73721414 0.73555094 0.73298683  
0.7339386 0.7382355 0.73872063 0.73601774]  
average cross\_validation score: 0.74  
k: 11  
Cross-validation score:  
[0.73409563 0.73520444 0.72598753 0.73513514 0.73374913 0.73215523  
0.73338416 0.73698801 0.73608705 0.73483956]  
average cross\_validation score: 0.73  
k: 12  
Cross-validation score:  
[0.73153153 0.73014553 0.72356202 0.73146223 0.73146223 0.72772003  
0.72957239 0.73456234 0.73428512 0.73359207]  
average cross\_validation score: 0.73  
k: 13  
Cross-validation score:  
[0.72806653 0.72910603 0.71898822 0.72626473 0.72543313 0.72460152  
0.72458244 0.72957239 0.72853282 0.72624576]  
average cross\_validation score: 0.73  
k: 14  
Cross-validation score:  
[0.72127512 0.72404712 0.71330561 0.72072072 0.71850312 0.71961192  
0.71869152 0.7225726 0.7225726 0.71876083]  
average cross\_validation score: 0.72

In [9]: *#Decision tree on test data*

```

tree_train = DecisionTreeClassifier(max_depth=6,random_state=2).fit(X_train, y_train)
y_predicted = tree_train.predict(X_test)
curTestAccuracy = tree_train.score(X_test, y_test)
print(curTestAccuracy)
print(metrics.classification_report(y_predicted, y_test))

```

0.7362366247158618

	precision	recall	f1-score	support
1	1.00	0.73	0.85	35277
2	0.00	0.00	0.00	0
3	0.00	0.00	0.00	1
4	0.01	0.57	0.02	68
5	0.00	0.00	0.00	0
6	1.00	0.90	0.95	728
7	0.00	0.00	0.00	0
avg / total	1.00	0.74	0.85	36074

/anaconda/envs/py35/lib/python3.5/site-packages/sklearn/metrics/classification.py:1137: Undefined 'recall', 'true', average, warn\_for)

In [71]: *# train ridge regression*

```

ridge_model = linear_model.Ridge(alpha = 0.3)
ridge_model.fit(X_train, y_train)
ridge_predicted = ridge_model.predict(X_test)

```

In [72]: *#ridge classifier*

```

ridge = RidgeClassifier().fit(X_train, y_train)
y_ridge = ridge.predict(X_test)
curTestAccuracy = ridge.score(X_test, y_test)
print(curTestAccuracy)
print(metrics.classification_report(y_ridge, y_test))

```

0.7379732739420936

	precision	recall	f1-score	support
1	1.00	0.73	0.85	44048
2	0.00	0.00	0.00	0
3	0.00	0.00	0.00	0
4	0.00	0.00	0.00	0
5	0.00	0.00	0.00	0
6	1.00	0.92	0.96	852
7	0.00	0.00	0.00	0

```
avg / total          1.00          0.74          0.85          44900
```

```
/anaconda/envs/py35/lib/python3.5/site-packages/sklearn/metrics/classification.py:1137: UndefinedLabelWarning:
  'recall', 'true', average, warn_for)
```

```
In [48]: #normalize the data
         from sklearn.preprocessing import MinMaxScaler
```

```

mms=MinMaxScaler()
X_train_norm = mms.fit_transform(X_train)
X_test_norm = mms.fit_transform(X_test)
```

```
In [49]: #Question 2b Euclidian
         from sklearn.neighbors import KNeighborsClassifier
         import matplotlib.pyplot as plt
         %matplotlib inline
```

```

def plotNumNeighborsAccuracy():
    training_accuracy=[]
    test_accuracy=[]
    neighbor_settings = range(1,7)
    for curKvalue in neighbor_settings:
        #Build the model
        clf = KNeighborsClassifier(n_neighbors=curKvalue)
        fold accuracies = cross_val_score(clf, X_train_norm, y_train, cv=kfold)
        print("Cross-validation score:\n{}".format(fold accuracies))
        print("average cross_validation score: {:.2f}".format(fold accuracies.mean()))
```

```
In [50]: plotNumNeighborsAccuracy()
```

```
Cross-validation score:
```

```
[0.56633081 0.59302814 0.59596808 0.58378832 0.57328853 0.58672827
 0.56404872 0.57034859 0.5800084  0.5749685 ]
```

```
average cross_validation score: 0.58
```

```
Cross-validation score:
```

```
[0.66708648 0.69046619 0.67744645 0.68584628 0.66190676 0.6900462
 0.65896682 0.68038639 0.67660647 0.67198656]
```

```
average cross_validation score: 0.68
```

```
Cross-validation score:
```

```
[0.67464316 0.68836623 0.67912642 0.68962621 0.66400672 0.69088618
 0.66274675 0.68962621 0.67576648 0.67786644]
```

```
average cross_validation score: 0.68
```

```
Cross-validation score:
```

```
[0.68178002 0.70558589 0.68542629 0.69802604 0.67366653 0.70222596
 0.67198656 0.70096598 0.69088618 0.69088618]
```

```
average cross_validation score: 0.69
```



```

Cross-validation score:
[0.70025189 0.71146577 0.69424612 0.70180596 0.68416632 0.71608568
 0.67576648 0.70306594 0.6900462 0.69256615]
average cross_validation score: 0.70
Cross-validation score:
[0.7031906 0.71650567 0.70138597 0.70852583 0.68710626 0.72280554
 0.68164637 0.70558589 0.70096598 0.69970601]
average cross_validation score: 0.70

```

```

In [73]: # train polynomial feature model

```

```

poly_features = PolynomialFeatures(degree=2, include_bias=False)
X_poly = poly_features.fit_transform(X)
print(X.shape)
print(X_poly.shape)

X_polyTrain, X_polyTest, y_polyTrain, y_polyTest = train_test_split(X_poly, y, random_s
poly_model = linear_model.LinearRegression().fit(X_polyTrain, y_polyTrain)
y_polyPredicted = poly_model.predict(X_polyTest)
curTestAccuracy = poly_model.score(X_polyTest, y_polyTest)
print(curTestAccuracy)

```

```

(224499, 8)
(224499, 44)
0.15513881704609067

```

```

In [74]: # linear regression model

```

```

from sklearn import linear_model
lr_model = linear_model.LinearRegression().fit(X_train,y_train)
# predicted Vs actual values
y_predicted = lr_model.intercept_ + lr_model.coef_*X_test

```

```

In [108]: from sklearn.preprocessing import StandardScaler

```

```

stdsc = StandardScaler()
stdsc.fit(X_train)
x_train_std = stdsc.transform(X_train)
x_test_std = stdsc.transform(X_test)

```

```

In [105]: from sklearn.neural_network import MLPClassifier

```

```

mlp = MLPClassifier(random_state=42)
mlp.fit(x_train_std, y_train)
print("Accuracy on training set: {:.2f}".format(mlp.score(x_train_std, y_train)))
print("Accuracy on test set: {:.2f}".format(mlp.score(x_test_std, y_test)))

```

Accuracy on training set: 0.73  
Accuracy on test set: 0.74

```
In [106]: for i in range(1,11):
            mlp2 = MLPClassifier(activation='tanh', hidden_layer_sizes=[i])
            mlp2.fit(x_train_std, y_train)
            print("Accuracy on training set: {:.2f}".format(mlp2.score(x_train_std, y_train)))
            print("Accuracy on test set: {:.2f}".format(mlp2.score(x_test_std, y_test)))
```

Accuracy on training set: 0.73  
Accuracy on test set: 0.74  
Accuracy on training set: 0.73  
Accuracy on test set: 0.74  
Accuracy on training set: 0.73  
Accuracy on test set: 0.73  
Accuracy on training set: 0.73  
Accuracy on test set: 0.73  
Accuracy on training set: 0.73  
Accuracy on test set: 0.74  
Accuracy on training set: 0.73  
Accuracy on test set: 0.74  
Accuracy on training set: 0.73  
Accuracy on test set: 0.74  
Accuracy on training set: 0.73  
Accuracy on test set: 0.74

```
/anaconda/envs/py35/lib/python3.5/site-packages/sklearn/neural_network/multilayer_perceptron.py:
    warnings.warn("Training interrupted by user.")
```

Accuracy on training set: 0.73  
Accuracy on test set: 0.73  
Accuracy on training set: 0.73  
Accuracy on test set: 0.73

```
In [61]: import numpy as np
import pandas as pd
from sklearn.datasets import load_digits
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.cluster import KMeans
from sklearn.metrics import accuracy_score
from sklearn.svm import SVC
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score

class cluster():
```

```

def load(self,data):

    crimes = pd.factorize(data['offence_type'])
    df_2['offence_type'] = crimes[0]
    crime_list = crimes[1]

    location_type = pd.factorize(data['location_type'])
    data['location_type'] = location_type[0]
    location_type_list = location_type[1]

    X = data
    y = data['offence_type'].values

    self.X_train, self.X_test, self.y_train, self.y_test = train_test_split(X, y, t

def __init__(self, data):
    self.load(data)

def classify(self, model=LogisticRegression(random_state=42)):
    model.fit(self.X_train, self.y_train)
    y_pred = model.predict(self.X_test)

    print(classification_report(self.y_test,y_pred, target_names=crime_list))
    print('Accuracy: {}'.format(accuracy_score(self.y_test, y_pred)))

def Kmeans(self, output='include'):
    n_clusters = len(np.unique(self.y_train))
    clf = KMeans(n_clusters = n_clusters, random_state=42)
    clf.fit(self.X_train)
    y_labels_train = clf.labels_
    y_labels_test = clf.predict(self.X_test)
    if output == 'include':
        self.X_train['km_clust'] = y_labels_train
        self.X_test['km_clust'] = y_labels_test
    elif output == 'exclude':
        self.X_train = y_labels_train[:, np.newaxis]
        self.X_test = y_labels_test[:, np.newaxis]
    else:
        raise ValueError('output should be either add or replace')
    return self

```

```
In [38]: df = pd.read_csv('crimes.csv',sep=',')
```

```

columns = ['year', 'month', 'day', 'hour', 'location_type', 'zipcode', 'x_coord', 'y_coord', 'o

df_2 = df[columns]
df_2 = df_2.dropna(axis=0)
cluster(df_2).Kmeans(output='include').classify()

```

/anaconda/envs/py35/lib/python3.5/site-packages/ipykernel\_launcher.py:46: SettingWithCopyWarning  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#>  
/anaconda/envs/py35/lib/python3.5/site-packages/ipykernel\_launcher.py:47: SettingWithCopyWarning  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#>

	precision	recall	f1-score	support
Theft	0.73	1.00	0.85	48388
Auto Theft	0.00	0.00	0.00	1556
Aggravated Assault	0.00	0.00	0.00	3583
Burglary	0.00	0.00	0.00	3480
Robbery	0.00	0.00	0.00	9014
Rape	0.89	1.00	0.94	1149
Murder	0.00	0.00	0.00	47
avg / total	0.54	0.74	0.63	67217

Accuracy: 0.7364952318609875

/anaconda/envs/py35/lib/python3.5/site-packages/sklearn/metrics/classification.py:1135: Undefined  
'precision', 'predicted', average, warn\_for)

In [40]: cluster(df\_2).Kmeans(output='include').classify(model=RidgeClassifier())

	precision	recall	f1-score	support
Theft	0.90	1.00	0.95	48388
Auto Theft	0.00	0.00	0.00	1556
Aggravated Assault	0.00	0.00	0.00	3583
Burglary	0.00	0.00	0.00	3480
Robbery	0.72	1.00	0.84	9014
Rape	0.93	1.00	0.96	1149
Murder	0.00	0.00	0.00	47

avg / total	0.76	0.87	0.81	67217
-------------	------	------	------	-------

Accuracy: 0.870434562685035

```
/anaconda/envs/py35/lib/python3.5/site-packages/ipykernel_launcher.py:46: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/indexing.html#
/anaconda/envs/py35/lib/python3.5/site-packages/ipykernel_launcher.py:47: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/indexing.html#
/anaconda/envs/py35/lib/python3.5/site-packages/sklearn/metrics/classification.py:1135: UndefinedLabelWarning:
'precision', 'predicted', average, warn_for)
```